

Supporting Appendix: Condensed Version

The Foreign Policy Impact of Trade-Based Status Gains: When China Overtakes the US as
Top Export Destination

Prepared for Chinese Political Science Review

This file contains the condensed supporting appendix. The complete diagnostics, tables, prompt documentation, and robustness analyses are available in Online Resource 1, supplied as a separate anonymized PDF with the submission.

Included sections

- Brazil and China Ideal-Point Series
- Triple-Difference Pre-Trends
- Cross-Country Robustness: Duration Thresholds
- Cross-Country Sample Audit: Goods-Only Treatment Coding
- Public Cue and Recoverability Audit
- Measurement Robustness: UNGA-DM Ideal Points

Available in Online Resource 1

- Resolution-Level UNGA Vote Diagnostics
- SDiD
- Selected Public-Cue Cross-Country Analyses
- ChatGPT Classification
- Cross-Country: Goods-Only Treated Countries

Brazil and China Ideal-Point Series

Figure 7 plots the separate UNGA ideal-point series for Brazil and China. This diagnostic clarifies the interpretation of the absolute-distance outcome used in the main SDiD design: China's average position changes little around 2009 despite sizable year-to-year volatility, while Brazil shows a sustained post-2009 movement toward China.




Fig. 7 Brazil and China UNGA ideal-point series. The dashed line marks 2009. The figure separates the two series behind the absolute-distance outcome used in the main SDiD analysis. Note: UNGA ideal-point score (Bailey et al. scale). Window: 1990-2022 annual series.

Triple-Difference Pre-Trends

Figure 8 reports the input to the identification assumption behind the triple difference. Each series is Brazil's average distance to China minus its distance to the United States, net of the donor-pool average in the same year, computed separately for human-rights votes and for every other vote where China and the United States voted differently. What the triple difference requires is that these two series move together before treatment, so that any common drift cancels when one domain is differenced against the other.




Fig. 8 Brazil's distance from the donor pool in China-US divergent votes, by issue domain. Each point is Brazil's mean vote-level distance to China minus its distance to the United States, minus the donor-pool mean in the same year; more negative values mean Brazil sits closer to China than the donors do. The dashed line marks 2009. The two domains track each other before 2009 and separate afterward, which is the pattern the triple difference assumes and then estimates. Note: Vote-level distance to China minus distance to the United States, relative to the donor-pool mean. Window: 2005-2012 roll-call votes where China and the United States voted differently.

Cross-Country Robustness: Duration Thresholds

The main cross-country estimate uses a five-year durability threshold because the theory concerns status that persists long enough to become politically usable. Table 18 reports the same goods-only IFE design under three- and seven-year thresholds, along with the clean single-entry robustness sample. The estimates remain negative across the restricted-risk-set rows, though precision varies with treated-unit support.

Table 18 Duration-threshold robustness for goods-only China top export-destination status.

Minimum duration	Specification	ATT	SE	95% CI	p-value	Units (T/C)	Treated periods	r*
3	Status-current, restricted risk set	-0.113	0.041	[-0.193, -0.034]	0.005	39/126	477	2
3	Clean single-entry robustness	-0.121	0.041	[-0.202, -0.040]	0.003	26/126	332	2
5	Status-current, restricted risk set	-0.101	0.039	[-0.177, -0.024]	0.010	35/126	440	2
5	Clean single-entry robustness	-0.122	0.041	[-0.202, -0.042]	0.003	25/126	329	2
7	Status-current, restricted risk set	-0.094	0.040	[-0.173, -0.016]	0.019	31/126	411	2
7	Clean single-entry robustness	-0.122	0.042	[-0.205, -0.039]	0.004	23/126	318	2

Note:

ATT in absolute UNGA ideal-point distance to China; lower values indicate convergence toward China. SE: bootstrap with 10,000 replications.

Window: goods-only ITPD-E rank metric. The restricted-risk-set rows retain clean pre-entry non-China-top years and qualifying treated years.

The clean single-entry rows additionally keep only treated countries with one observed qualifying China-top period.

Cross-Country Sample Audit: Goods-Only Treatment Coding

The revised rank metric excludes services before partner ranks are computed. Table 19 documents the ITPD-E broad-sector filter. Table 20 then summarizes how the five-year goods-only rule classifies countries in the UNGA/trade panel.

Table 19 ITPD-E sector filter used to construct the goods-only China top export-destination metric.

Broad sector	Raw rows	Positive-trade rows	Included in rank metric
Agriculture	7331458	2484549	Yes
Manufacturing	71208330	28203707	Yes
Mining and Energy	1205032	416413	Yes
Services	1002558	519520	No

Table 20 Country-level sample audit for the five-year goods-only China top export-destination rule.

Role	Countries
Excluded: no clean observed prior non-China-top year	1
Excluded: pre-2000 China-top entry	2
Excluded: short China-top periods	26
Never-China-top controls	128
Qualifying treated countries	35

Note:

The audit uses the goods-only rank metric and classifies countries before the estimator-specific screen for minimum untreated support. Country counts are therefore a coding audit, not a replacement for the treated/control counts in the estimation table.

Public Cue and Recoverability Audit

To assess how closely the empirical treatment used in the panel regression – China becoming a country’s largest destination for goods exports – maps onto the theorized treatment of a publicly salient status change, I use a source audit originally constructed for 14 countries in an earlier no-covariate treated sample. This audit is not a census of the 35 countries treated in the current goods-only restricted-risk-set specification: 13 audited countries remain in the current treated sample, Qatar does not, and 22 current treated countries were not included in the source audit. The underlying matrix records both recovered and unrecovered China cues. Table 21 displays the nine cases in which the China cue is currently coded as unknown; Chile, Brazil, Uruguay, Saudi Arabia, and Qatar have recovered cues and are omitted by the display filter.

The table distinguishes weak observation from cases in which public rank language about the displaced incumbent is recoverable. For the China cue, `medium` means one countable strong or moderate first-window source with explicit rank language. For the former-incumbent benchmark, `medium` means either one countable source with explicit incumbent-rank or displacement language, or at least two countable trade-coverage records from at least one independent source without a clear top-rank label. A countable source must be eligible, fall within one year of the source-audit entry year, be coded strong or moderate, enter the benchmark, and not be marked `DO_NOT_COUNT`. `Unknown` means that the incumbent was not identified or that no countable benchmark source was recovered; no case is currently coded `low`. Non-recovery is not evidence that a cue was absent because contemporaneous local or official records may be inaccessible or inconsistent with the trade-rank series. Except for Australia, whose displayed year has been reconciled to the current 2009 entry, the table retains the entry windows used in the legacy source audit. Table 22 reports Australia’s broad-trade source separately.

Table 21: Selected unresolved cases from a 14-country legacy source audit of public rank-language recoverability and metric caveats.

Country	Entry year	China cue	Ex-#1 benchmark	Interpretation	Caveat
---------	------------	-----------	-----------------	----------------	--------

Solomon Islands	2003	unknown	unknown	Weak observation/recoverability problem	Additional local/official archive work required.
Philippines	2005	unknown	unknown	Weak observation/recoverability problem	Treatment-rank audit issue in contemporaneous sources.
Angola	2007	unknown	unknown	Weak observation/recoverability problem	Additional local/official archive work required.
Malaysia	2009	unknown	medium	Benchmark recoverable; follow up before recoding salience	Rank-definition and China/Hong Kong aggregation caveat.
Sierra Leone	2012	unknown	unknown	Weak observation/recoverability problem	Local statistics conflict with Belgium-as-incumbent.
Myanmar (Burma)	2014	unknown	unknown	Weak observation/recoverability problem	Additional local/official archive work required.
Saudi Arabia	2013	unknown	unknown	Weak observation/recoverability problem	Corrected 2013-2014 public-cue window remains unresolved; rank-definition reconciliation is required.
Kuwait	2018	unknown	unknown	Weak observation/recoverability problem	Additional local/official archive work required.

Note:

Nine unresolved cases from a 14-country legacy audit; this is not a census of the current treated sample. Definitions and the display filter reported in the text. Australia's entry year follows the current goods-only annual panel; other entry windows remain those of the legacy audit.

Table 22 Australian Financial Review context source for Australia.

Source	Australian Financial Review
Date	2009-11-09
Metric	Australia's largest aggregate trading partner
Counted in export-destination benchmark	No
Raw file	data/raw/ex_top1_salience/AUS/2009/ australian_financial_review/ aus_afr_china_trading_partner_2009.html
Interpretation	Counts for broad public-cue salience in the current 2009 treatment window; remains excluded from the strict export-destination benchmark because the metric combines exports and imports of goods and services.

Note:

The AFR item documents a broad aggregate-trading-partner cue. It documents broad-metric salience during the 2009 annual treatment window. The publication date does not establish that a strict goods-export cue preceded treatment.

Measurement Robustness: UNGA-DM Ideal Points

The main estimations in the paper used the classic ideal-point measurement by BSV (Bailey et al. 2017) as the response variable. However, according to Fjelstul et al. (2026), the set of resolutions included in the BSV dataset is not complete. They provided a broader dataset, along with re-estimations of ideal points based on the new data. To assess whether the results of the paper change with this new measurement, I reran the main estimations with this new response variable. To make the comparison as close as possible, I held the treatment definition, donor and risk-set rules, covariate structure, and inference conventions fixed. The UNGA-DM series has a shorter time window, and the SDiD unit and time weights and the IFE factor count are re-estimated by their algorithms in response to the outcome series. These induced adjustments mean that the comparison evaluates the robustness of the analysis package to the alternative outcome; it is not a pure measurement-only perturbation of an otherwise identical fitted counterfactual. It is also important to note that 28 country-session rows had no ISO3c match and 0 SDiD rows had no outcome in the alternative measurement.

The two outcome measures preserve the negative direction of the Brazilian SDiD estimate, but they do not support a blanket claim of robustness across the Brazilian and cross-country designs. Table 25 shows the comparison. Under UNGA-DM, the point estimate is nominally higher than the original in absolute terms (-0.331 versus -0.273), with a lower p-value of 0.008. The directional rank is the same in both measurements (3/96 (0.031)). The bilateral rank improves from 7/96 (0.073) to 5/96 (0.052). The pre-treatment means differ (0.647 under BSV versus 0.837 under UNGA-DM). But the effect measured as a percentage of the pre-treatment average is similar: -42.2 and -39.5 percent. The Brazilian result therefore retains its direction and inferential support under the alternative measure, with a similar relative magnitude. The cross-country comparison below is more sensitive to the outcome source: its IFE point estimate is closer to zero on the UNGA-DM scale and its uncertainty interval includes zero. The 2×2 comparison below holds the common window and country-year panel fixed while varying the outcome source and factor count, allowing conditional factor-count comparisons within each measure. It does not provide a unique, scale-invariant decomposition of the difference between outcomes.

Table 25 Brazil SDiD under two sources of UNGA ideal points. Both columns use the same panel, the same donor pool, and the same placebo standard error settings; the outcome series is the only input that differs, and the synthetic-control weights are re-estimated from it.

	Bailey, Strezhnev and Voeten	Fjelstul, Hug and Kilby
ATT	-0.273	-0.331
Placebo standard error	0.131	0.125
95 percent confidence interval	[-0.529, -0.017]	[-0.576, -0.085]
p (normal approximation, two-sided)	0.037	0.008
Directional placebo rank (p)	3/96 (0.031)	3/96 (0.031)
Two-sided absolute placebo rank (p)	7/96 (0.073)	5/96 (0.052)
Intercept-adjusted pre-treatment RMSPE	0.058	0.058
Brazil pre-treatment outcome mean	0.647	0.837
ATT as percent of pre-treatment mean	-42.2	-39.5

Note:

Estimates are the post-2009 gap in absolute UNGA ideal-point distance to China, 1997-2015 annual panel, no covariates. Placebo SEs use 20,000 replications and seed 20260520 in both columns. Placebo-in-space ranks involve no random draws and are exactly reproducible.

Figure 13 reports the placebo-in-space distribution for the UNGA-DM SDiD fit. It is close to the distribution obtained under the original outcome. The directional placebo rank is unchanged under the alternative measure (3/96 (0.031) in both columns), while the bilateral absolute rank improves from 7/96 (0.073) to 5/96 (0.052). Table 25 reports this cross-measure comparison, and Table 26 reports the donor-exclusion checks under UNGA-DM.




Fig. 13 Placebo-in-space distribution under the UNGA-DM outcome, with Brazil marked. Brazil's estimate is the 3rd most negative of 96 assignments (directional $p = 0.031$); the two-sided absolute rank is 5/96 ($p = 0.052$). Negative values indicate convergence toward China. Note: Estimated post-2009 gap in absolute UNGA ideal-point distance to China, UNGA-DM measure. Window: 1997-2015 annual SDiD panel.

Table 26 Placebo rank inference under the UNGA-DM outcome, by donor-exclusion criterion.

Comparison set	Excluded donors	Directional rank (p)	Two-sided rank (p)
All valid assignments	–	3/96 (0.031)	5/96 (0.052)
Exclude goods-only China-top donor assignments (harmonized: window criterion)	–	3/96 (0.031)	5/96 (0.052)
Exclude goods-only China-top donor assignments (legacy 5-year qualifying criterion; audit only)	GAB, KWT	3/94 (0.032)	5/94 (0.053)

Note:

The window criterion excludes donors observed as China-top in goods exports within 1997-2015, which is the criterion used elsewhere in the paper. The legacy criterion excludes countries that qualify under the five-year duration rule outside the analysis window, and is reported for audit only.

Donor weights are estimated separately for each fit, so they are not identical across outcomes. Table 27 reports how much they change. Weight is spread among donors in a similar way in both fits: 0.044 under UNGA-DM and 0.031 under BSV. The correlation of the donor pool weights is moderate (0.462), but among the top ten donors, 6 are the same. The donor mix therefore shifts moderately across outcomes. For Brazil, the negative estimate and similar relative magnitude are preserved, but the two fits do not use identical counterfactual weights; this comparison does not establish stability of the cross-country result.

Table 27 Ten largest synthetic-control donor weights under the UNGA-DM fit, with the weight the same donor receives under the main outcome.

Donor	UNGA-DM weight	BSV weight
COL	0.044	0.022
IND	0.031	0.021
SEN	0.031	0.009
BRB	0.029	0.015
GTM	0.027	0.027
GHA	0.026	0.000
TTO	0.026	0.015
PRY	0.026	0.027
BOL	0.025	0.023
SGP	0.023	0.029

Note:

Weights are normalized over the full donor pool of 95 units rather than over the displayed donors. The ten donors shown carry 28.8 percent of total weight under the UNGA-DM fit. Across all donors, the correlation between the two weight vectors is 0.462, and 6 of the ten largest donors are shared by the two fits.

Regarding the panel regression, Table 28 shows the IFE estimation in the same time window for both outcomes, and the original estimation with the larger time window. The estimate retains the negative sign but is closer to zero on the UNGA-DM scale, with an uncertainty interval that includes zero: -0.027 with $p = 0.573$. However, the models are not the same, since the number of latent factors chosen by cross-validation is different: 2 under

BSV and 1 under UNGA-DM. The 2×2 comparison below holds the common window and country-year panel fixed while crossing the outcome source with one versus two latent factors, allowing sensitivity to factor count to be assessed within each outcome measure rather than uniquely decomposing the cross-outcome difference.

Table 28 Cross-country interactive fixed-effects panel under the two outcome sources, in a common window with identical rows.

Specification	ATT	SE	95 percent CI	p	Factors (CV)	Country-years
BSV full window 1990-2022 (paper main, stored target)	-0.101	0.039	[-0.177, -0.024]	0.010	2	5,037
BSV common window (≤ 2020 , identical rows)	-0.095	0.040	[-0.173, -0.018]	0.016	2	4,727
UNGA-DM common window (≤ 2020 , identical rows)	-0.027	0.048	[-0.120, 0.066]	0.573	1	4,727

Note:

All rows use the goods-only status-current restricted risk set with 35 treated and 126 control countries. The common-window rows keep exactly the same country-years in both outcome sources. Confidence intervals come from 10,000 bootstrap replications. ATTs and intervals are expressed on each outcome's original scale; raw cross-outcome differences do not identify a scale-invariant measurement contribution.

Table 29 holds the common window and country-year panel fixed and crosses the two outcome sources with one and two latent factors. This 2×2 comparison describes sensitivity to factor count conditional on the outcome measure. Holding the outcome fixed, moving from one to two factors makes the ATT more negative (larger in absolute value): from -0.036 to -0.095 under BSV, and from -0.027 to -0.065 under UNGA-DM. Cross-validation selects 2 factors under BSV and 1 under UNGA-DM. The change from one to two factors differs across the two measures, indicating a descriptive interaction between outcome source and factor count rather than a unique additive decomposition. This comparison does not test the interaction. Both UNGA-DM fits have negative point estimates and uncertainty intervals that include zero; overlapping intervals do not establish equivalence. Holding the factor count at two still leaves a smaller UNGA-DM estimate (-0.065) than the BSV estimate (-0.095), but

these raw estimates are expressed on different outcome scales. Their difference is a conditional sensitivity comparison, not a scale-invariant contribution attributable to outcome source.

Table 29 Outcome source by number of latent factors, holding the panel fixed.

Outcome	Factors (fixed)	Selected by CV	ATT	SE	95 percent CI	p
BSV	1	No	-0.036	0.037	[-0.109, 0.037]	0.336
UNGA-DM	1	Yes	-0.027	0.048	[-0.120, 0.066]	0.573
BSV	2	Yes	-0.095	0.040	[-0.173, -0.018]	0.016
UNGA-DM	2	No	-0.065	0.055	[-0.172, 0.043]	0.238

Note:

Same common window and identical rows as the previous table. Paired bootstrap on UNGA-DM minus BSV in their original outcome units, 1,000 replications and no failures: 0.069 (percentile $p = 0.134$) with the original CV-selected counts held fixed (BSV: two; UNGA-DM: one), and 0.030 (percentile $p = 0.356$) with two factors for both outcomes. The same stratified country draws are used for both measures; factor counts are not reselected within replications. These are scale-dependent contrasts, not a unique decomposition of factor and outcome-source components. Non-rejection does not establish equivalence.

To sum up, the Brazilian SDiD result retains its negative direction and inferential support across the two outcome measurements, with similar relative magnitude but a different absolute scale. The 2×2 IFE comparison shows that two factors produce a more negative ATT under either outcome, with a different within-measure change under BSV and UNGA-DM. Both UNGA-DM fits have negative point estimates and uncertainty intervals that include zero; this does not establish equivalence between them. Even with two factors held fixed, the raw UNGA-DM estimate is closer to zero than the BSV estimate on its own scale. This conditional comparison does not isolate a scale-invariant outcome-source contribution beyond factor selection. The paired test does not detect a statistically significant difference between the ATTs, but this non-rejection does not establish equivalence and may reflect uncertainty.